

Session 2: Guided Exercise

Handwritten Digit Recognition with Feedforward Neural Networks

Neural Networks Course - Computer Engineering

May 9, 2025

Abstract

In this guided exercise, you will implement a complete feedforward neural network for recognizing handwritten digits from the MNIST dataset using PyTorch. You will learn how to load and preprocess data, define a neural network architecture, train the model, and evaluate its performance. This exercise will reinforce theoretical concepts from the lecture and give you hands-on experience with neural network implementation.

Contents

1	Introduction	2
2	Setting Up the Environment	2
3	MNIST Dataset Overview	2
4	Loading and Preprocessing Data	3
5	Defining the Neural Network Architecture	4
6	Loss Function and Optimizer	5
7	Training the Model	5
8	Evaluating the Model	7
9	Visualizing Results	7
9.1	Training Metrics	7
9.2	Visualizing Predictions	8
10	Saving and Loading the Model	8
11	Extending the Exercise	9
12	Comparison with State-of-the-Art	10
13	Conclusion	10
14	References	11
15	Appendix: Complete Code	11

1 Introduction

The MNIST dataset (Modified National Institute of Standards and Technology) is a large collection of handwritten digits commonly used for training and testing machine learning algorithms. It contains 70,000 grayscale images of handwritten digits (0-9), each of size 28×28 pixels.

In this guided exercise, you will build a feedforward neural network (also known as a multilayer perceptron) to recognize these handwritten digits. By the end of this exercise, you will have implemented a complete deep learning pipeline and learned how to:

- Load and preprocess image data
- Define a neural network architecture
- Train a model using backpropagation and gradient descent
- Evaluate model performance
- Visualize results and model predictions

Note

This exercise assumes basic familiarity with Python and PyTorch. If you are new to PyTorch, please refer to the documentation at <https://pytorch.org/docs/stable/index.html>.

2 Setting Up the Environment

Before starting, ensure that you have all the necessary libraries installed. You will need PyTorch, torchvision, matplotlib, and numpy. If you are using Google Colab, these libraries are pre-installed.

```
1 pip install torch torchvision matplotlib numpy
```

Listing 1: Installing required packages

Alternatively, you can use the provided Jupyter notebook that already has all dependencies configured.

Tip

If you encounter CUDA-related errors and don't have a GPU, you can force CPU usage with: `device = torch.device("cpu")`

3 MNIST Dataset Overview



Figure 1: Sample images from the MNIST dataset

The MNIST dataset has the following characteristics:

- 60,000 training images and 10,000 test images
- 10 classes (digits 0-9)
- Grayscale images with size 28×28 pixels
- Pixel values range from 0 (white) to 255 (black)

Figure 1 shows sample images from the dataset. Notice the variations in handwriting styles, which make this a good challenge for classification algorithms.

4 Loading and Preprocessing Data

The first step is to load the MNIST dataset and prepare it for training. PyTorch provides convenient utilities through the `torchvision.datasets` module.

```

1 import torch
2 from torchvision import datasets, transforms
3 from torch.utils.data import DataLoader
4
5 # Define transformations
6 transform = transforms.Compose([
7     transforms.ToTensor(), # Convert images to PyTorch tensors
8     transforms.Normalize((0.1307,), (0.3081,)) # Normalize with mean and std of
9     MNIST
10 ])
11
12 # Download and load training data
13 train_dataset = datasets.MNIST(root='./data',
14                                train=True,
15                                download=True,
16                                transform=transform)
17
18 # Download and load test data
19 test_dataset = datasets.MNIST(root='./data',
20                               train=False,
21                               download=True,
22                               transform=transform)
23
24 # Create data loaders for batch processing
25 batch_size = 64
26 train_loader = DataLoader(train_dataset,
27                           batch_size=batch_size,
28                           shuffle=True)
29
30 test_loader = DataLoader(test_dataset,
31                          batch_size=batch_size,
32                          shuffle=False)

```

Listing 2: Loading the MNIST dataset

Note

The `Normalize` transform standardizes the pixel values using the mean (0.1307) and standard deviation (0.3081) of the MNIST dataset. This helps in faster convergence during training.

5 Defining the Neural Network Architecture

Now, let's define our feedforward neural network architecture. We'll create a simple network with one hidden layer.




Figure 2: Architecture of our feedforward neural network

```
1 import torch.nn as nn
2 import torch.nn.functional as F
3
4 class FeedforwardNeuralNet(nn.Module):
5     """A simple feedforward neural network with one hidden layer"""
6
7     def __init__(self, input_size, hidden_size, output_size):
8         """Initialize the network architecture
9
10        Args:
11            input_size: Number of input features
12            hidden_size: Number of neurons in the hidden layer
13            output_size: Number of output classes
14        """
15        super(FeedforwardNeuralNet, self).__init__()
16
17        # First fully connected layer (input  $\rightarrow$  hidden)
18        self.fc1 = nn.Linear(input_size, hidden_size)
19
20        # Second fully connected layer (hidden  $\rightarrow$  output)
21        self.fc2 = nn.Linear(hidden_size, output_size)
22
23        # Dropout layer for regularization (preventing overfitting)
24        self.dropout = nn.Dropout(0.2)
25
26    def forward(self, x):
27        """Forward pass through the network
28
29        Args:
30            x: Input tensor of shape [batch_size, 1, 28, 28]
31
32        Returns:
33            Output tensor of shape [batch_size, output_size]
34        """
35        # Reshape input: [batch_size, 1, 28, 28]  $\rightarrow$  [batch_size, 784]
36        x = x.view(-1, 28*28)
37
38        # First layer with ReLU activation
39        x = F.relu(self.fc1(x))
40
41        # Apply dropout
42        x = self.dropout(x)
43
44        # Output layer (no activation yet - will use softmax with loss function)
```

```

45         x = self.fc2(x)
46
47         return x
48
49 # Initialize the model
50 input_size = 28 * 28 # MNIST images are 28x28 pixels
51 hidden_size = 128    # Number of neurons in the hidden layer
52 output_size = 10     # 10 digits (0-9)
53 model = FeedforwardNeuralNet(input_size, hidden_size, output_size)

```

Listing 3: Defining the neural network model

Tip

The choice of hidden layer size (128 neurons) is a hyperparameter. You can experiment with different values to see how it affects model performance.

6 Loss Function and Optimizer

To train our model, we need to define:

- A loss function to measure how well the model is performing
- An optimizer to update the model parameters based on the gradients

```

1 import torch.optim as optim
2
3 # Define the loss function (Cross-Entropy Loss)
4 criterion = nn.CrossEntropyLoss()
5
6 # Define the optimizer (Stochastic Gradient Descent)
7 learning_rate = 0.01
8 optimizer = optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9)

```

Listing 4: Setting up loss function and optimizer

Note

The Cross-Entropy Loss combines softmax activation with negative log-likelihood loss, making it suitable for multi-class classification problems.

7 Training the Model

Now, let's implement the training loop. We'll train the model for multiple epochs (complete passes through the training dataset).

```

1 # Training parameters
2 num_epochs = 5
3
4 # Lists to store metrics for plotting
5 train_losses = []
6 train_accuracies = []
7
8 def train(model, train_loader, criterion, optimizer, epoch):
9     """Train the model for one epoch"""
10    # Set model to training mode
11    model.train()
12

```

```

13     running_loss = 0.0
14     correct = 0
15     total = 0
16
17     # Iterate over batches
18     for batch_idx, (data, target) in enumerate(train_loader):
19         # Clear gradients from previous step
20         optimizer.zero_grad()
21
22         # Forward pass
23         outputs = model(data)
24
25         # Calculate loss
26         loss = criterion(outputs, target)
27
28         # Backward pass
29         loss.backward()
30
31         # Update weights
32         optimizer.step()
33
34         # Accumulate loss
35         running_loss += loss.item()
36
37         # Calculate accuracy
38         _, predicted = torch.max(outputs.data, 1)
39         total += target.size(0)
40         correct += (predicted == target).sum().item()
41
42         # Print statistics every 100 batches
43         if (batch_idx + 1) % 100 == 0:
44             print(f'Epoch [{epoch+1}/{num_epochs}], '
45                   f'Step [{batch_idx+1}/{len(train_loader)}], '
46                   f'Loss: {loss.item():.4f}, '
47                   f'Accuracy: {100 * correct / total:.2f}%')
48
49         # Calculate average metrics
50         epoch_loss = running_loss / len(train_loader)
51         epoch_acc = 100 * correct / total
52
53         # Store for plotting
54         train_losses.append(epoch_loss)
55         train_accuracies.append(epoch_acc)
56
57     return epoch_loss, epoch_acc
58
59 # Train for multiple epochs
60 for epoch in range(num_epochs):
61     epoch_loss, epoch_acc = train(model, train_loader, criterion, optimizer,
62                                   epoch)
63     print(f'Epoch {epoch+1}/{num_epochs} completed - '
64           f'Loss: {epoch_loss:.4f}, Accuracy: {epoch_acc:.2f}%')

```

Listing 5: Training loop

Warning

Training deep learning models can be computationally intensive. If you're running this on a CPU, it might take several minutes to complete.

8 Evaluating the Model

After training, we need to evaluate how well our model performs on unseen data (the test set).

```
1 def evaluate(model, test_loader):
2     """Evaluate model performance on the test set"""
3     # Set model to evaluation mode
4     model.eval()
5
6     correct = 0
7     total = 0
8
9     # Disable gradient calculation
10    with torch.no_grad():
11        for data, target in test_loader:
12            # Forward pass
13            outputs = model(data)
14
15            # Get predictions
16            _, predicted = torch.max(outputs.data, 1)
17
18            # Update statistics
19            total += target.size(0)
20            correct += (predicted == target).sum().item()
21
22    # Calculate and return accuracy
23    accuracy = 100 * correct / total
24    print(f'Test Accuracy: {accuracy:.2f}%')
25
26    return accuracy
27
28 # Evaluate on test set
29 test_accuracy = evaluate(model, test_loader)
```

Listing 6: Model evaluation

Note

We set the model to evaluation mode (`model.eval()`) during testing. This disables dropout and batch normalization layers, which behave differently during training and evaluation.

9 Visualizing Results

Let's visualize the training progress and model predictions to better understand how our model performs.

9.1 Training Metrics

First, let's plot the loss and accuracy during training:

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(12, 4))
4
5 plt.subplot(1, 2, 1)
6 plt.plot(train_losses)
7 plt.title('Training Loss')
8 plt.xlabel('Epoch')
9 plt.ylabel('Loss')
10
```

```

11 plt.subplot(1, 2, 2)
12 plt.plot(train_accuracies)
13 plt.title('Training Accuracy')
14 plt.xlabel('Epoch')
15 plt.ylabel('Accuracy (%)')
16
17 plt.tight_layout()
18 plt.savefig('training_metrics.png')
19 plt.show()

```

Listing 7: Plotting training metrics

Figure 3: Training loss and accuracy over epochs

9.2 Visualizing Predictions

Now, let's visualize how well our model predicts on the test set:

```

1 def visualize_predictions(model, test_loader, num_samples=10):
2     """Visualize model predictions on sample test images"""
3     # Set model to evaluation mode
4     model.eval()
5
6     # Get a batch of test data
7     examples = iter(test_loader)
8     samples, labels = next(examples)
9
10    # Make predictions
11    with torch.no_grad():
12        outputs = model(samples)
13        _, predicted = torch.max(outputs, 1)
14
15    # Plot results
16    plt.figure(figsize=(15, 3))
17    for i in range(num_samples):
18        plt.subplot(1, num_samples, i+1)
19        plt.imshow(samples[i][0], cmap='gray')
20
21        # Green title for correct predictions, red for incorrect
22        if predicted[i] == labels[i]:
23            plt.title(f'Pred: {predicted[i]}\nTrue: {labels[i]}', color='green')
24        else:
25            plt.title(f'Pred: {predicted[i]}\nTrue: {labels[i]}', color='red')
26
27        plt.axis('off')
28
29    plt.tight_layout()
30    plt.savefig('model_predictions.png')
31    plt.show()
32
33 # Visualize predictions
34 visualize_predictions(model, test_loader)

```

Listing 8: Visualizing model predictions

10 Saving and Loading the Model

Once you've trained a model that performs well, you'll want to save it for future use.



Figure 4: Model predictions on test images (green: correct, red: incorrect)

```

1 # Save the model
2 torch.save(model.state_dict(), 'mnist_feedforward_model.pth')
3 print("Model saved successfully!")
4
5 # Later, to load the model:
6 loaded_model = FeedforwardNeuralNet(input_size, hidden_size, output_size)
7 loaded_model.load_state_dict(torch.load('mnist_feedforward_model.pth'))
8 loaded_model.eval() # Set to evaluation mode

```

Listing 9: Saving and loading the model

11 Extending the Exercise

Now that you have implemented a basic neural network for MNIST classification, you can extend your learning by trying the following challenges:

Challenge

1. Add more hidden layers to create a deeper network
2. Experiment with different activation functions (Sigmoid, Tanh, Leaky ReLU)
3. Implement learning rate scheduling to improve convergence
4. Try different optimizers (Adam, RMSprop)
5. Add batch normalization for faster and more stable training
6. Implement early stopping to prevent overfitting
7. Use data augmentation to improve generalization
8. Convert your model to a Convolutional Neural Network (CNN)

Here's an example of how to create a deeper network:

```

1 class DeepFeedforwardNet(nn.Module):
2     """A deeper feedforward neural network with multiple hidden layers"""
3
4     def __init__(self, input_size, hidden_sizes, output_size):
5         """Initialize the network architecture
6
7         Args:
8             input_size: Number of input features
9             hidden_sizes: List of sizes for each hidden layer
10            output_size: Number of output classes
11        """
12        super(DeepFeedforwardNet, self).__init__()
13
14        # Create a list to hold all layers

```

```

15     layers = []
16
17     # Input layer
18     layers.append(nn.Linear(input_size, hidden_sizes[0]))
19     layers.append(nn.ReLU())
20     layers.append(nn.BatchNorm1d(hidden_sizes[0]))
21     layers.append(nn.Dropout(0.2))
22
23     # Hidden layers
24     for i in range(len(hidden_sizes) - 1):
25         layers.append(nn.Linear(hidden_sizes[i], hidden_sizes[i+1]))
26         layers.append(nn.ReLU())
27         layers.append(nn.BatchNorm1d(hidden_sizes[i+1]))
28         layers.append(nn.Dropout(0.2))
29
30     # Output layer
31     layers.append(nn.Linear(hidden_sizes[-1], output_size))
32
33     # Combine all layers into a sequential model
34     self.model = nn.Sequential(*layers)
35
36     def forward(self, x):
37         """Forward pass through the network"""
38         x = x.view(-1, 28*28) # Flatten the input
39         return self.model(x)
40
41 # Example usage:
42 deeper_model = DeepFeedforwardNet(
43     input_size=28*28,
44     hidden_sizes=[256, 128, 64],
45     output_size=10
46 )

```

Listing 10: Implementing a deeper network

12 Comparison with State-of-the-Art

The feedforward neural network we implemented achieves reasonable performance on MNIST (typically around 97-98% accuracy), but it's worth noting that more advanced architectures like Convolutional Neural Networks (CNNs) can achieve over 99.5% accuracy on this dataset.

Model Architecture	Typical Accuracy on MNIST
Feedforward Neural Network (1 hidden layer)	~97-98%
Deep Feedforward Network (3+ hidden layers)	~98-99%
Convolutional Neural Network (CNN)	~99-99.5%
CNN with data augmentation	>99.5%

Table 1: Comparison of different architectures on MNIST classification

13 Conclusion

In this guided exercise, you implemented a complete feedforward neural network for handwritten digit recognition using the MNIST dataset. You learned how to:

- Load and preprocess image data
- Define a neural network architecture

- Implement the training and evaluation loops
- Visualize training progress and model predictions
- Save and load trained models

This implementation serves as a foundation for more complex neural network architectures and applications. The concepts and techniques you’ve learned are directly applicable to other image classification tasks and can be extended to more challenging datasets.

14 References

1. LeCun, Y., Cortes, C., & Burges, C. (2010). MNIST handwritten digit database. Retrieved from <http://yann.lecun.com/exdb/mnist/>
2. Paszke, A., et al. (2019). PyTorch: An Imperative Style, High-Performance Deep Learning Library. Advances in Neural Information Processing Systems 32 (NeurIPS 2019).
3. Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press. <http://www.deeplearningbook.org>

15 Appendix: Complete Code

The complete code for this exercise is available in the accompanying Python file `mnist_feedforward.py` and can also be found in the course GitHub repository. You can run it directly or use it as a reference for your own implementation.